

CMSC5702 (Advanced Topics in Parallel / Distributed Systems)

Assignment 1: Parallel Programming with Sobel Filter

Objective

Apply the Sobel filter for edge detection on **large square grayscale images** using MPI and OpenMP, and compare the performance of shared-memory vs distributed-memory parallelization. Optionally, include a blurring step to reduce noise before applying the Sobel operator.

This assignment allows you to practice:

- Sequential programming
 - OpenMP parallelization on shared-memory systems
 - MPI parallelization on distributed-memory systems
 - Performance measurement, verification, and speedup analysis
-

Problem Description

The **Sobel filter** computes approximate image gradients to detect edges. For a grayscale image (I):

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} * I, \quad G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} * I$$

The **gradient magnitude** at each pixel is:

$$G = \sqrt{G_x^2 + G_y^2}$$

Optional **blurring step**: compute the local average of pixels in a small neighborhood (e.g., 3x3) before applying the Sobel operator to reduce noise.

Part 1: Sequential Version

- Implement a standard sequential C program to apply the Sobel operator on a grayscale image.
- If including blurring, apply an average filter first (e.g., 3x3 kernel) before Sobel.

Question: Explain the algorithm you implemented and if you applied blurring, explain the blurring method.

Your Answer:

My implementation follows a two-stage process.

First, I implemented the optional **blurring step** to reduce image noise before edge detection. The method I used is the 3×3 **average filter**, which is also called a mean filter. As shown in my `blur_filter` function, this filter works by sliding a 3x3 window over the image. For each pixel, it calculates the average value of all 9 pixels in its 3×3 neighborhood (summing them up and dividing by 9.0) and uses this average as the new value for that pixel.

Second, I apply the **Sobel operator** to the *blurred* image. This is implemented in my `sobel_filter` function. This algorithm calculates two gradients for each pixel by convolving the image with two 3x3 kernels:

- G_x : Using the 3x3 Sobel kernel to detect horizontal changes.
- G_y : Using the 3x3 Sobel kernel to detect vertical changes.

Finally, the magnitude of the gradient is computed using the formula $G = \sqrt{G_x^2 + G_y^2}$ to get the final pixel intensity for the edge-detected image. In my implementation, for pixels on the 1-pixel border of the image, I set their output value to 0.0, as the 3x3 kernel cannot be fully applied there.

```
In [1]: !pip install pandas
```

```
Requirement already satisfied: pandas in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (2.2.3)
Requirement already satisfied: numpy>=1.26.0 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from pandas) (2.0.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from pandas) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from pandas) (2024.1)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
```

```
In [2]: !pip install matplotlib
```

```
Requirement already satisfied: matplotlib in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (3.9.2)
Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (4.55.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (1.4.7)
Requirement already satisfied: numpy>=1.23 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (2.0.1)
Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=8 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/hwpytorch/lib/python3.12/site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
```

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

seq_df = pd.read_csv("sequential_times.csv")
```

```
In [4]: seq_df
```

```
Out[4]:   image_size  run1_time  run2_time  run3_time
0          256    0.000933    0.001047    0.000888
1         1024    0.015544    0.014636    0.014826
2         4000    0.210384    0.194488    0.215198
3        16000    4.079864    3.984604    3.932855
```

```
In [5]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

# Assuming your data is in a DataFrame called df
# Calculate average and standard deviation
seq_df['average_time'] = seq_df[['run1_time', 'run2_time', 'run3_time']].mean()
seq_df['std_dev'] = seq_df[['run1_time', 'run2_time', 'run3_time']].std(axis=1)

# Plot with error bars showing standard deviation
plt.figure(figsize=(8,5))
plt.errorbar(seq_df['image_size'], seq_df['average_time'],
             yerr=seq_df['std_dev'],
             marker='o',
             color='blue',
```

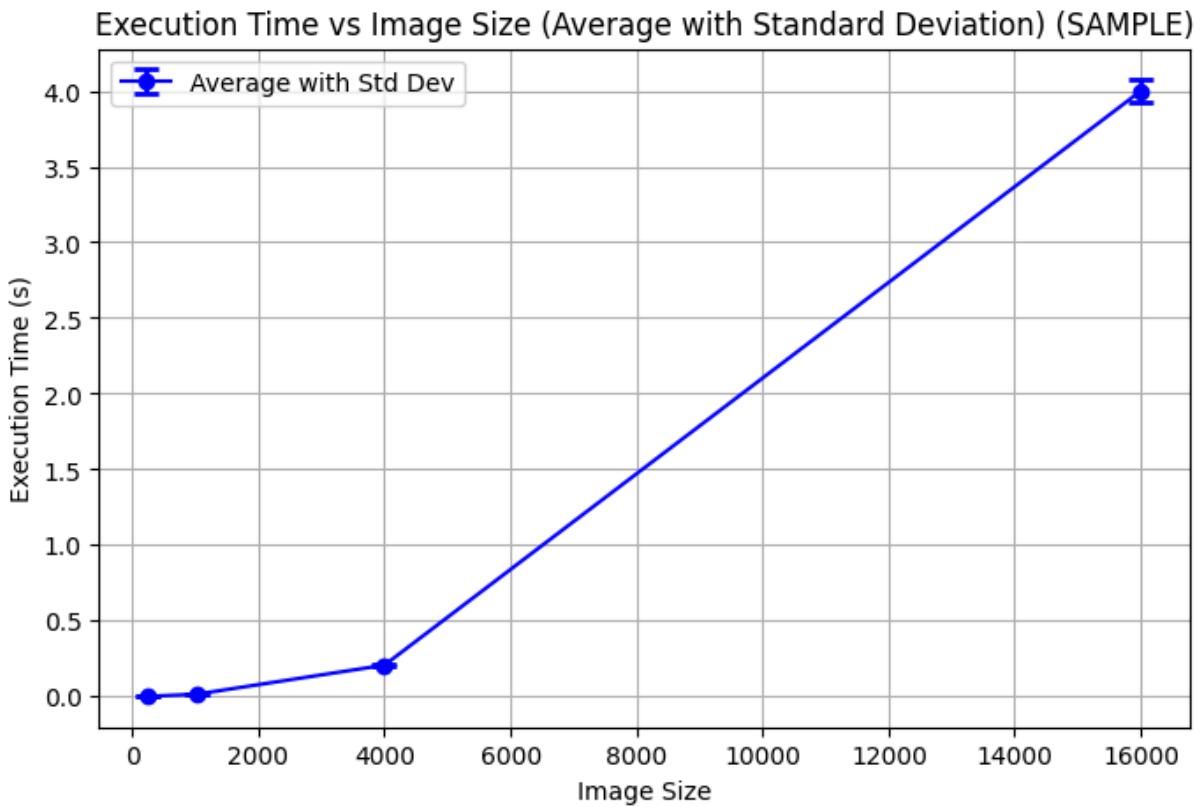
```

        label='Average with Std Dev',
        capsize=5, # Adds caps to the error bars
        capthick=2,
        elinewidth=2)

plt.xlabel('Image Size')
plt.ylabel('Execution Time (s)')
plt.title('Execution Time vs Image Size (Average with Standard Deviation) (S')
plt.legend()
plt.grid(True)
plt.show()

# Optional: Print the calculated values
print("Data with averages and standard deviations:")
print(seq_df[['image_size', 'average_time', 'std_dev']])

```



Data with averages and standard deviations:

	image_size	average_time	std_dev
0	256	0.000956	0.000082
1	1024	0.015002	0.000479
2	4000	0.206690	0.010838
3	16000	3.999108	0.074570

Part 2: OpenMP Programming

- Use `#pragma omp parallel for collapse(2)` for nested loops.
- Experiment with `schedule(static)` and other scheduling clauses.
- Process large images in tiles (e.g., 512×512 or 1024×1024) to reduce memory bandwidth issues.

- [Bonus Point] Optionally include a blurring step before the Sobel operator to reduce noise.

Question: Explain the algorithm you implemented and if you applied blurring, explain the blurring method.

Your Answer:

For the OpenMP implementation, my algorithm follows the same two-stage process as the sequential version, but the parallelization strategy is optimized for cache performance.

I implemented the **blurring step**. The method is the 3×3 **average filter**, where each pixel's new value is the average of its 9-pixel neighborhood (summing them up and dividing by 9.0) to reduce noise.

The core of my parallel implementation is a **tiling (or block-based) decomposition**. Here is the logic:

1. **Tiling:** Instead of parallelizing the image row by row, I divide the entire image into small, fixed-size square blocks.
2. **Parallelism:** I use OpenMP's `#pragma omp parallel for collapse(2)` directive. This directive assigns *entire tiles* to different threads. The `collapse(2)` tells OpenMP to merge the two outer loops (the one iterating tile rows and the one iterating tile columns) into one large loop, which provides many independent tasks (the tiles) for the threads to work on.

This tiling method is applied to both functions:

- First, `blur_filter_tiled` runs, with threads processing different tiles to create the blurred image.
- Second, `sobel_filter_tiled` runs on the blurred image, using the exact same parallel tiling strategy. Inside each tile, the standard Sobel logic is applied (convolving with Gx and Gy kernels, then calculating the magnitude

$$G = \sqrt{G_x^2 + G_y^2}.$$

Processing the image in tiles ensures that the data a thread is working on (a single tile) is small enough to fit well within the CPU caches, which reduces memory bandwidth issues and improves performance, especially on large images.

```
In [6]: omp_df = pd.read_csv("openmp_times.csv")
```

```
In [7]: omp_df.head()
```

Out[7]:

	image_size	threads	run1_time	run2_time	run3_time
0	256	1	0.002602	0.002662	0.002512
1	256	2	0.002633	0.002648	0.002686
2	256	4	0.002697	0.002705	0.002729
3	256	8	0.002934	0.002855	0.002830
4	256	16	0.003115	0.003054	0.003121

In [8]:

```

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

# Load CSV if not already loaded
# omp_df = pd.read_csv("openmp_timings.csv")

# ① Compute average and standard deviation across runs
omp_df['average_time'] = omp_df[['run1_time','run2_time','run3_time']].mean()
omp_df['std_dev'] = omp_df[['run1_time','run2_time','run3_time']].std(axis=1)

# Prepare columns for speedups
omp_df['absolute_speedup'] = 0.0
omp_df['relative_speedup'] = 0.0

# ② Compute absolute and relative speedup
# Absolute speedup: vs sequential execution (threads=1)
# Relative speedup: vs OpenMP 1-thread execution (same in this dataset)
for size in omp_df['image_size'].unique():
    # Sequential time (threads=1)
    seq_time = omp_df[(omp_df['image_size']==size) & (omp_df['threads']==1)]

    subset_idx = omp_df[omp_df['image_size']==size].index
    omp_df.loc[subset_idx, 'absolute_speedup'] = seq_time / omp_df.loc[subset_idx, 'run1_time']
    omp_df.loc[subset_idx, 'relative_speedup'] = seq_time / omp_df.loc[subset_idx, 'run1_time']

# ③ Plot runtime with error bars
plt.figure(figsize=(8,5))
for size in sorted(omp_df['image_size'].unique()):
    subset = omp_df[omp_df['image_size']==size]
    plt.errorbar(subset['threads'], subset['average_time'],
                 yerr=subset['std_dev'],
                 marker='o', label=f'{size}x{size}', capsized=5)

    plt.xlabel('Number of Threads')
    plt.ylabel('Runtime (s)')
    plt.title('OpenMP Runtime vs Threads (Average ± Std Dev) (SAMPLE)')
    plt.xscale('log', base=2)
    plt.xticks([1,2,4,8,16,32])
    plt.grid(True, which='both', linestyle='--', alpha=0.6)
    plt.legend()
    plt.show()

# ④ Plot absolute speedup

```

```

plt.figure(figsize=(8,5))
for size in sorted(omp_df['image_size'].unique()):
    subset = omp_df[omp_df['image_size']==size]
    plt.plot(subset['threads'], subset['absolute_speedup'], marker='o', label=f'{size}x')

plt.xlabel('Number of Threads')
plt.ylabel('Absolute Speedup')
plt.title('OpenMP Absolute Speedup vs Threads (SAMPLE)')
plt.xscale('log', base=2)
plt.xticks([1,2,4,8,16,32])
plt.grid(True, which='both', linestyle='--', alpha=0.6)
plt.legend()
plt.show()


# ⑤ Plot relative speedup (same as absolute here)
plt.figure(figsize=(8,5))
for size in sorted(omp_df['image_size'].unique()):
    subset = omp_df[omp_df['image_size']==size]
    plt.plot(subset['threads'], subset['relative_speedup'], marker='o', label=f'{size}x')

plt.xlabel('Number of Threads')
plt.ylabel('Relative Speedup')
plt.title('OpenMP Relative Speedup vs Threads (SAMPLE)')
plt.xscale('log', base=2)
plt.xticks([1,2,4,8,16,32])
plt.grid(True, which='both', linestyle='--', alpha=0.6)
plt.legend()
plt.show()


# ⑥ Optional: Print the full dataframe
print("OpenMP data with runtime, std_dev, absolute and relative speedups:")
print(omp_df[['image_size','threads','average_time','std_dev','absolute_speedup','relative_speedup']])

```


OpenMP Runtime vs Threads (Average \pm Std Dev) (SAMPLE)



OpenMP Absolute Speedup vs Threads (SAMPLE)



OpenMP Relative Speedup vs Threads (SAMPLE)



OpenMP data with runtime, std_dev, absolute and relative speedups:

	image_size	threads	average_time	std_dev	absolute_speedup	\
0	256	1	0.002592	0.000075	1.000000	
1	256	2	0.002656	0.000027	0.976026	
2	256	4	0.002710	0.000017	0.956340	
3	256	8	0.002873	0.000054	0.902193	
4	256	16	0.003097	0.000037	0.837029	
5	256	32	0.006793	0.000634	0.381551	
6	1024	1	0.045112	0.001055	1.000000	
7	1024	2	0.023913	0.002435	1.886519	
8	1024	4	0.014483	0.000557	3.114776	
9	1024	8	0.008809	0.000231	5.121358	
10	1024	16	0.007126	0.000275	6.330963	
11	1024	32	0.007880	0.000090	5.724915	
12	4000	1	0.545054	0.034005	1.000000	
13	4000	2	0.284085	0.021516	1.918630	
14	4000	4	0.152469	0.007123	3.574859	
15	4000	8	0.086843	0.005820	6.276291	
16	4000	16	0.052449	0.002981	10.392076	
17	4000	32	0.043478	0.001073	12.536413	
18	16000	1	11.111987	0.696742	1.000000	
19	16000	2	6.314521	0.390359	1.759751	
20	16000	4	2.775402	0.044862	4.003739	
21	16000	8	1.384710	0.007167	8.024778	
22	16000	16	0.714140	0.014279	15.559964	
23	16000	32	0.528825	0.008575	21.012585	

	relative_speedup
0	1.000000
1	0.976026
2	0.956340
3	0.902193
4	0.837029
5	0.381551
6	1.000000
7	1.886519
8	3.114776
9	5.121358
10	6.330963
11	5.724915
12	1.000000
13	1.918630
14	3.574859
15	6.276291
16	10.392076
17	12.536413
18	1.000000
19	1.759751
20	4.003739
21	8.024778
22	15.559964
23	21.012585

Part 3: MPI Programming

- Choose a domain decomposition strategy, such as row-wise decomposition, column-wise decomposition, etc., decomposition. Your choice will affect performance, and better-designed decomposition and communication patterns will receive higher marks.
- For each process, you may need to include some approaches to allow seamless computation for the edges.
- [Bonus Point] You can include If a parallelized blurring step and you may need to try to optimize the communications.
- Verify your implementation using small test patterns (5×5 or 10×10) and visually inspect sample images for correctness.

```
In [9]: mpi_df = pd.read_csv("mpi_times.csv")
```

```
In [10]: mpi_df.head()
```

	image_size	nodes	processes	run1_time	run2_time	run3_time
0	256	1	1	0.004858	0.004847	0.004859
1	256	1	2	0.004317	0.004334	0.004342
2	256	1	4	0.002223	0.002965	0.002210
3	256	4	8	0.010011	0.009881	0.010153
4	256	4	16	0.010022	0.009914	0.010408

```
In [11]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

# Compute average and standard deviation of runtime
mpi_df['average_time'] = mpi_df[['run1_time', 'run2_time', 'run3_time']].mean()
mpi_df['std_dev'] = mpi_df[['run1_time', 'run2_time', 'run3_time']].std(axis=1)

# --- Compute Absolute and Relative Speedup from CSV ---

# Absolute speedup: sequential = 1 process on 1 node
absolute_speedup = {}
for img_size in mpi_df['image_size'].unique():
    # Find sequential runtime: 1 node, 1 process
    seq_time = mpi_df[(mpi_df['image_size']==img_size) &
                      (mpi_df['nodes']==1) &
                      (mpi_df['processes']==1)]['average_time'].values[0]
    absolute_speedup[img_size] = seq_time

# Add speedup columns
mpi_df['absolute_speedup'] = mpi_df.apply(lambda row: absolute_speedup[row['image_size']], axis=1)

# Relative speedup: 1 process on 1 node runtime as reference
relative_speedup = {}
```

```

for img_size in mpi_df['image_size'].unique():
    ref_time = mpi_df[(mpi_df['image_size']==img_size) &
                      (mpi_df['nodes']==1) &
                      (mpi_df['processes']==1)]['average_time'].values[0]
    relative_speedup[img_size] = ref_time

mpi_df['relative_speedup'] = mpi_df.apply(lambda row: relative_speedup[row['image_size']], axis=1)

# --- Plots ---

# Runtime with error bars
plt.figure(figsize=(10,6))
for img_size in sorted(mpi_df['image_size'].unique()):
    df_img = mpi_df[mpi_df['image_size']==img_size]
    plt.errorbar(df_img['processes'], df_img['average_time'], yerr=df_img['std_dev'], marker='o', capsized=5, label=f'{img_size}x{img_size}')
plt.xlabel('Number of Processes')
plt.ylabel('Average Runtime (s)')
plt.title('MPI: Average Runtime vs Number of Processes (SAMPLE)')
plt.legend(title='Image Size')
plt.grid(True)
plt.show()


# Absolute speedup
plt.figure(figsize=(10,6))
for img_size in sorted(mpi_df['image_size'].unique()):
    df_img = mpi_df[mpi_df['image_size']==img_size]
    plt.plot(df_img['processes'], df_img['absolute_speedup'], marker='o', label=f'{img_size}x{img_size}')
plt.xlabel('Number of Processes')
plt.ylabel('Absolute Speedup')
plt.title('MPI: Absolute Speedup vs Number of Processes (SAMPLE)')
plt.legend(title='Image Size')
plt.grid(True)
plt.show()

# Relative speedup
plt.figure(figsize=(10,6))
for img_size in sorted(mpi_df['image_size'].unique()):
    df_img = mpi_df[mpi_df['image_size']==img_size]
    plt.plot(df_img['processes'], df_img['relative_speedup'], marker='o', label=f'{img_size}x{img_size}')
plt.xlabel('Number of Processes')
plt.ylabel('Relative Speedup')
plt.title('MPI: Relative Speedup vs Number of Processes')
plt.legend(title='Image Size')
plt.grid(True)
plt.show()


# Optional: print dataframe
print(mpi_df[['image_size', 'nodes', 'processes', 'average_time', 'std_dev', 'absolute_speedup', 'relative_speedup']])

```


MPI: Average Runtime vs Number of Processes (SAMPLE)



MPI: Absolute Speedup vs Number of Processes (SAMPLE)



MPI: Relative Speedup vs Number of Processes



	image_size	nodes	processes	average_time	std_dev	absolute_speedup
0	256	1	1	0.004855	0.000007	1.000000
1	256	1	2	0.004331	0.000013	1.120911
2	256	1	4	0.002466	0.000432	1.968640
3	256	4	8	0.010015	0.000136	0.484740
4	256	4	16	0.010115	0.000260	0.479963
5	256	4	32	0.001689	0.000492	2.874852
6	1024	1	1	0.073846	0.001352	1.000000
7	1024	1	2	0.066292	0.001806	1.113961
8	1024	1	4	0.034436	0.000432	2.144431
9	1024	4	8	0.019822	0.000114	3.725473
10	1024	4	16	0.010837	0.000087	6.814069
11	1024	4	32	0.005906	0.000102	12.504318
12	4000	1	1	1.095608	0.007461	1.000000
13	4000	1	2	0.979258	0.000564	1.118815
14	4000	1	4	0.519954	0.013461	2.107125
15	4000	4	8	0.253195	0.005454	4.327131
16	4000	4	16	0.131151	0.001393	8.353769
17	4000	4	32	0.068871	0.001714	15.908118
18	16000	1	1	19.388482	0.104251	1.000000
19	16000	1	2	16.949818	0.223479	1.143876
20	16000	1	4	8.740856	0.008950	2.218144
21	16000	4	8	4.196811	0.094911	4.619813
22	16000	4	16	2.194286	0.025026	8.835897
23	16000	4	32	1.132161	0.026099	17.125203
	relative_speedup					
0						1.000000
1						1.120911
2						1.968640
3						0.484740
4						0.479963
5						2.874852
6						1.000000
7						1.113961
8						2.144431
9						3.725473
10						6.814069
11						12.504318
12						1.000000
13						1.118815
14						2.107125
15						4.327131
16						8.353769
17						15.908118
18						1.000000
19						1.143876
20						2.218144
21						4.619813
22						8.835897
23						17.125203

In []: